


Mathematical simplification


- Neurons are either on or off: represented by binary value.



- Inputs from n neurons: x_i
- then added $\sum_i x_i w_i$
- get multiplied by weights at “synapses”: $x_i w_i$
- if above threshold, then neuron is “on”.
- Remark: add x_0 to the input vector, in order to write the bias as $b = x_0 w_0$ to get more compact form **$\mathbf{w}\mathbf{x}$** , rather than **$\mathbf{w}\mathbf{x} + b$**


Transfer function

- Step function:



- Used by Pitts & McCulloch (1942), and in Rosenblatt's Perceptron (1957)


- Sigmoid (used in many “modern” feed forward neural nets):



- Sigmoidal function is differentiable (good for deriving gradient decent learning rule; Backpropagation algorithm, Werbos 1974)

Neural networks (1940s/50s)

- Pitts and McCulloch (1942/3): “Formal” (mathematical) neurons thought of as processing units
- *Networks* can emulate any logical function.
- D. Hebb: Connections between neurons can change. “Learning rule”: strengthen proportional to correlation between activity of pre- and post synaptic neuron.



Perceptron (Rosenblatt, 1957)

- Probably the first “learning machine” that was actually built. Artificial neuron that solved a classification task (supervised learning):
- Given: N input vectors \mathbf{x} together with labels l (these labels are the “teaching” signal).
- Goal: for any given input, the output of the classifier should be the same as the label.



Perceptron learning


- Simplify: assume labels are binary, either -1 or +1.
➡ *Binary classification.*
- *adjust weights according to the correctness of the output until all input data in training set are classified correctly.*
- **error measure:** compare output of the neuron (y) to desired output (= label, l): both are either -1 or 1, so:
 - $y \cdot l = 1$: correct classification, then: $l - y = 0$
 - $y \cdot l = -1$: incorrect classification, then: $l - y = 2l$
- adjust weight vector \mathbf{w} for each misclassified input \mathbf{x} , by adding $l \cdot \mathbf{x}$. This turns \mathbf{w} towards/away from \mathbf{x} if $l=1/-1$
- can do $\mathbf{w} \leftarrow \mathbf{w} + c(l-y)\mathbf{x}$ for all training examples \mathbf{x}
- c : step size parameter called “learning rate”.

Representational power

- the decision boundary of the perceptron is a line:


$$\mathbf{w}\mathbf{x} + b$$

- the perceptron learning algorithm *converges* if input data are **linearly separable**.
- Novikoff (1962): Perceptron Convergence Theorem; first *margin-based error bound* (in a way the “dawn” of statistical learning theory)




XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR




XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR




XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR




XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR




XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR




XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR



XOR problem

- Perceptron algorithm can not classify input data which can not be separated by a line.
- For example XOR



- ***A single artificial neuron*** can express the boolean functions AND, OR, and NOT, **but not XOR.**

Homework

- Implement perceptron learning algorithm: N Inputs \mathbf{x}_j & labels l_j
 - ▶ Initialize weight vector \mathbf{w}
 - ▶ While there exist misclassified examples:
 - ▶ Compute output $y_j = \theta(\mathbf{w}\mathbf{x}_j)$
 - ▶ For each example, update the weights: $\mathbf{w} \leftarrow \mathbf{w} + c(l_j - y_j)\mathbf{x}_j$
- Play around with the parameter (learning rate) and the input data, and verify for yourself what the Perceptron can and can not do
 - ▶ Make a movie of Perceptron converging, and one of Perceptron failing on the XOR.
- What else do you notice?
 - ▶ Is every solution the same? If not, are some “better” than others in some sense?